

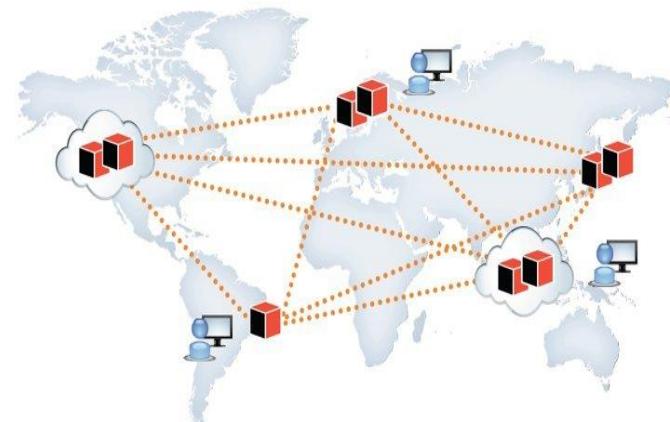
Undoing CRDT Operations Automatically

Provakar Mondal and Eli Tilevich
Software Innovations Lab, Virginia Tech, USA
`{provakar, tilevich}@cs.vt.edu`

Background

Modern distributed applications replicate data to different geo-locations to achieve:

- Low latency
- High availability
- Scalability
- Fault tolerance



Examples: Google Docs, Calendar, Figma, Facebook's OpenR, etc.

Background

- ❖ Distributed State Synchronization to ensure consistency
- ❖ Eventual Consistency
- ❖ Conflict-free Replicated Data Type (CRDT)

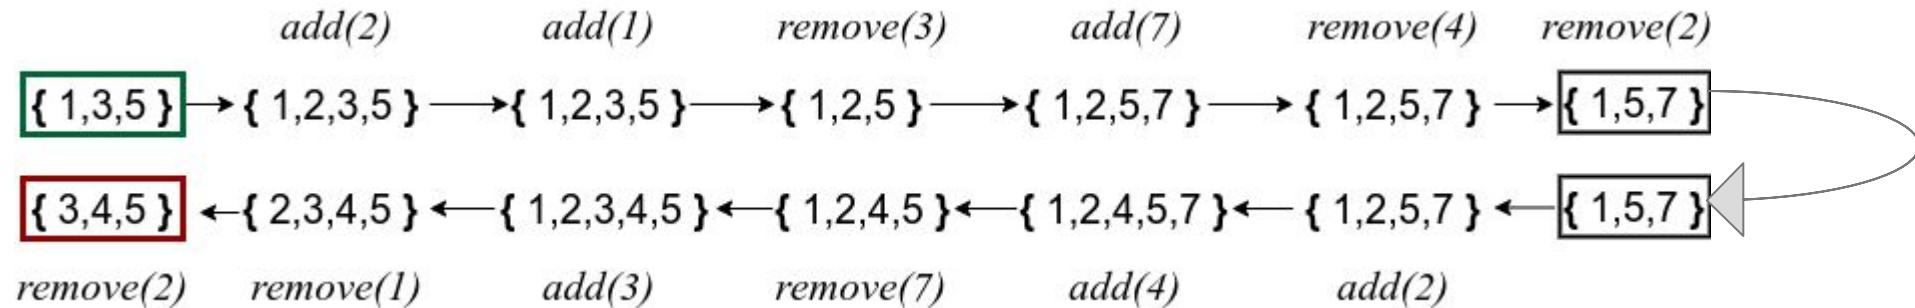
Problem Motivation

- Program errors cause *correct* updates of *incorrect* data
- CRDTs lack built-in *undo* functionality
- Ad-hoc undo implementations hard to reuse
- Modifying CRDT code by hand hurts portability

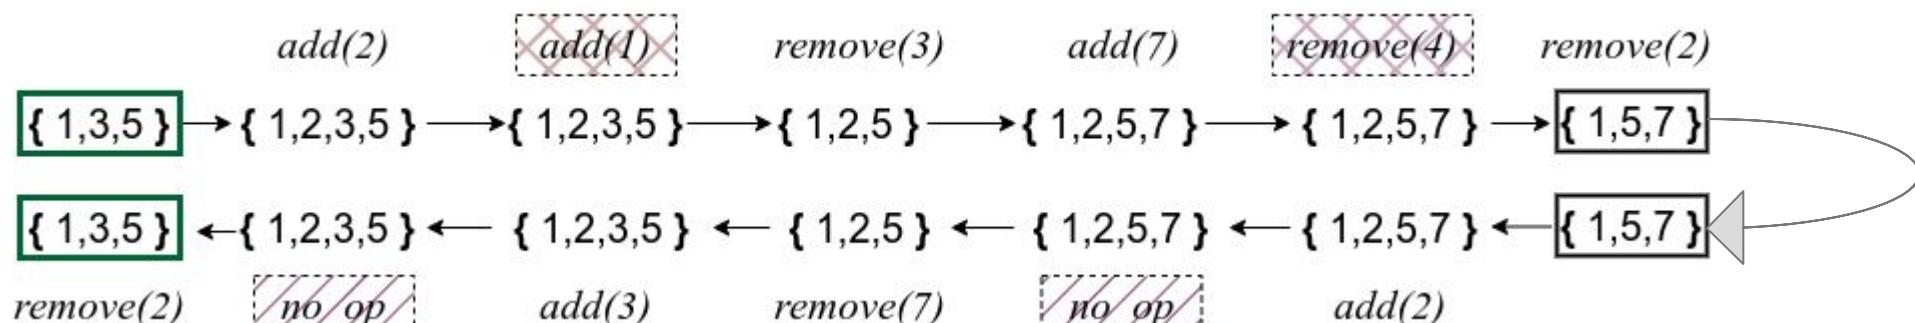
「 Paper's Contribution

- AUTO-UNDO
- Declarative metadata
 - undo pairs
 - no_op operation
- Undo Script
- No manual modifications of CRDT library

Motivating Example



Incorrect Undo Procedure



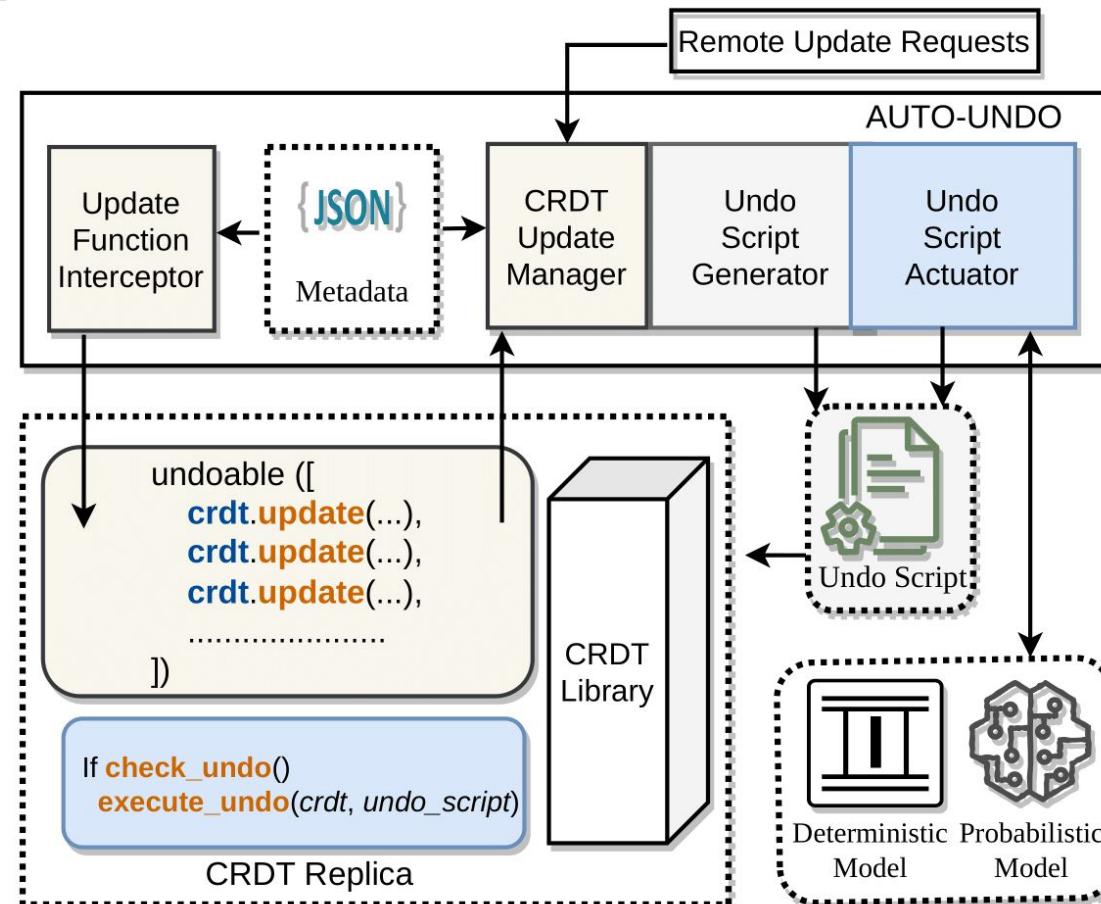
no_op Undo Procedure

AUTO-UNDO: Key Components

Four key components.

1. Update Function Interceptor
2. Update Manager
3. Undo Script Generator
4. Undo Script Actuator

Key Components of AUTO-UNDO

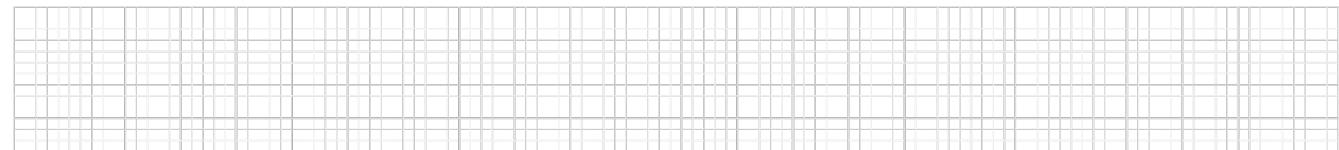


Undo Generation and Actuation with AUTO - UNDO

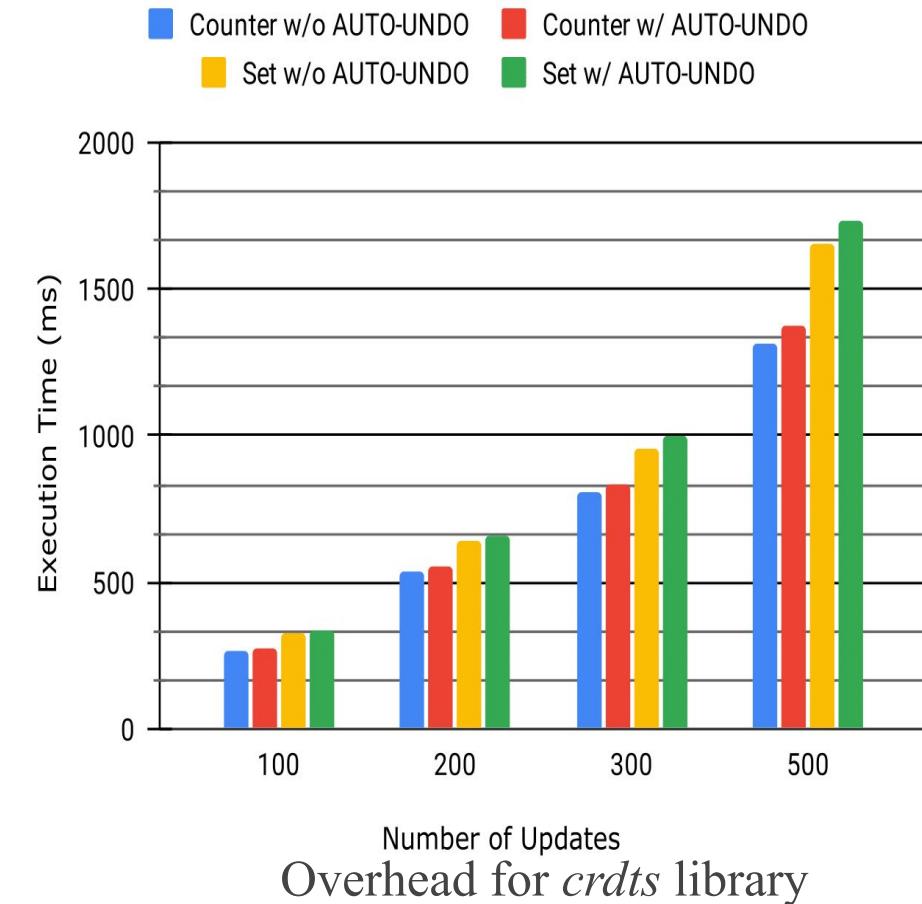
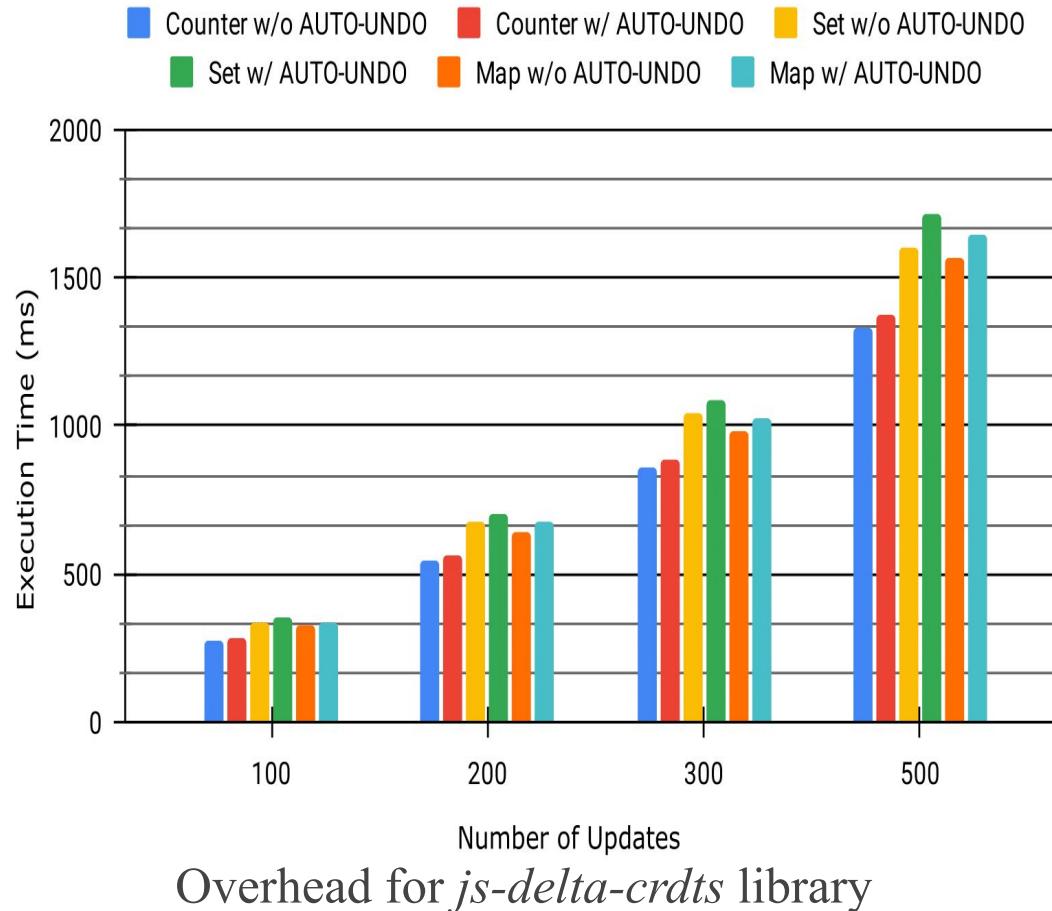
Evaluation

RQ1: What is the performance overhead of applying AUTO - UNDO to add the undo capability in existing CRDT libraries under different workloads?

RQ2: What is the accuracy/performance trade-offs between the deterministic and probabilistic models for determining whether to execute an undo procedure?



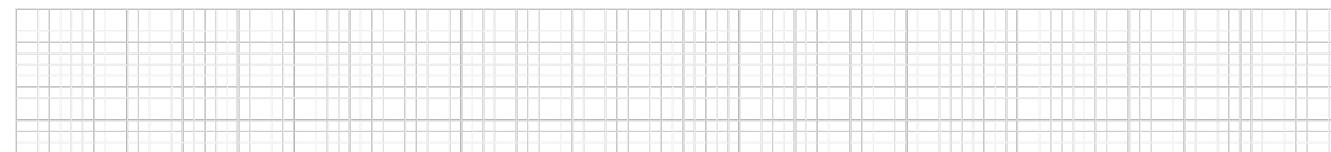
RQ1: Performance Overhead



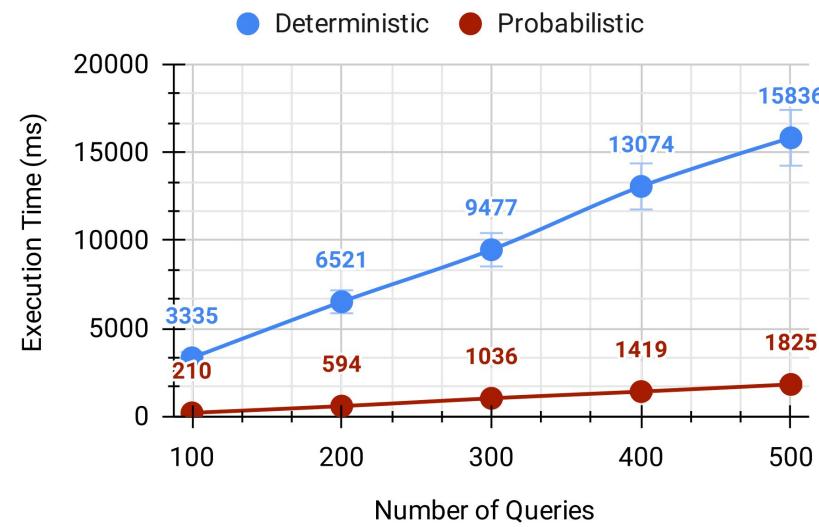
RQ1: Performance Overhead

| Number of Updates | <i>js-delta-crdts</i> | | | <i>crdts</i> | |
|----------------------|-----------------------|------|------|--------------|------|
| | Counter | Set | Map | Counter | Set |
| 100 | 2.93 | 3.87 | 3.36 | 2.64 | 3.08 |
| 200 | 2.19 | 3.98 | 4.18 | 3.13 | 3.45 |
| 300 | 3.27 | 4.32 | 4.37 | 3.49 | 4.07 |
| 500 | 3.61 | 4.94 | 4.78 | 4.33 | 4.56 |

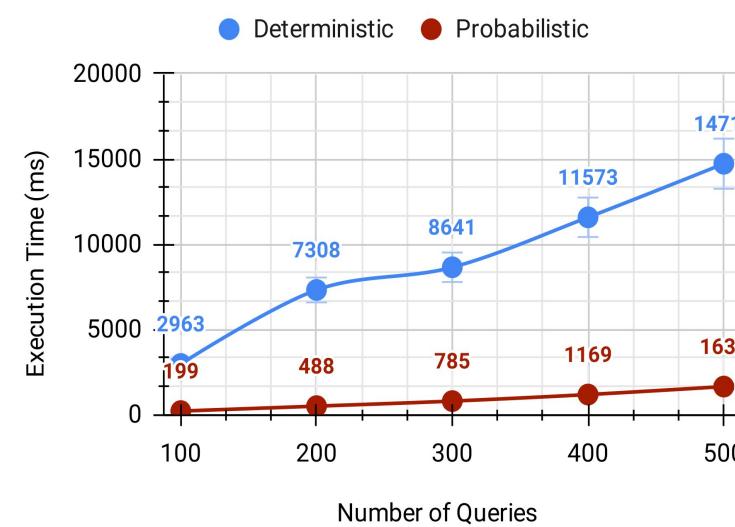
Latency Overhead (%)



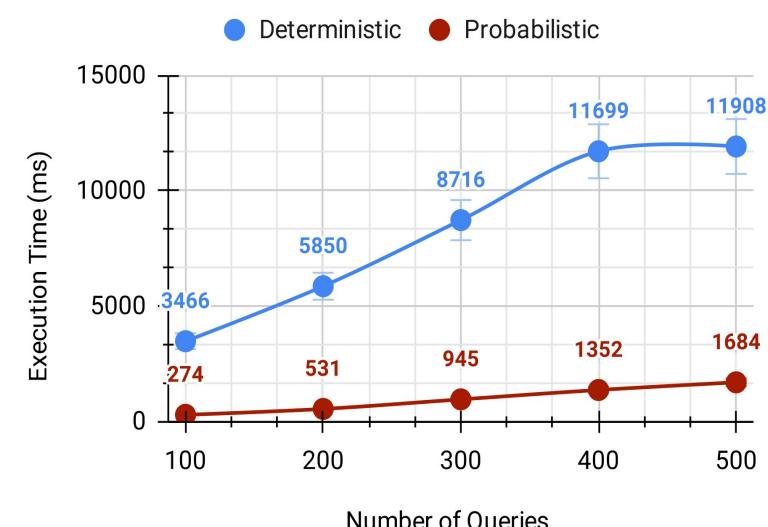
RQ2: Accuracy/Performance Trade-off



Add-Remove Set (js-delta-crdts)



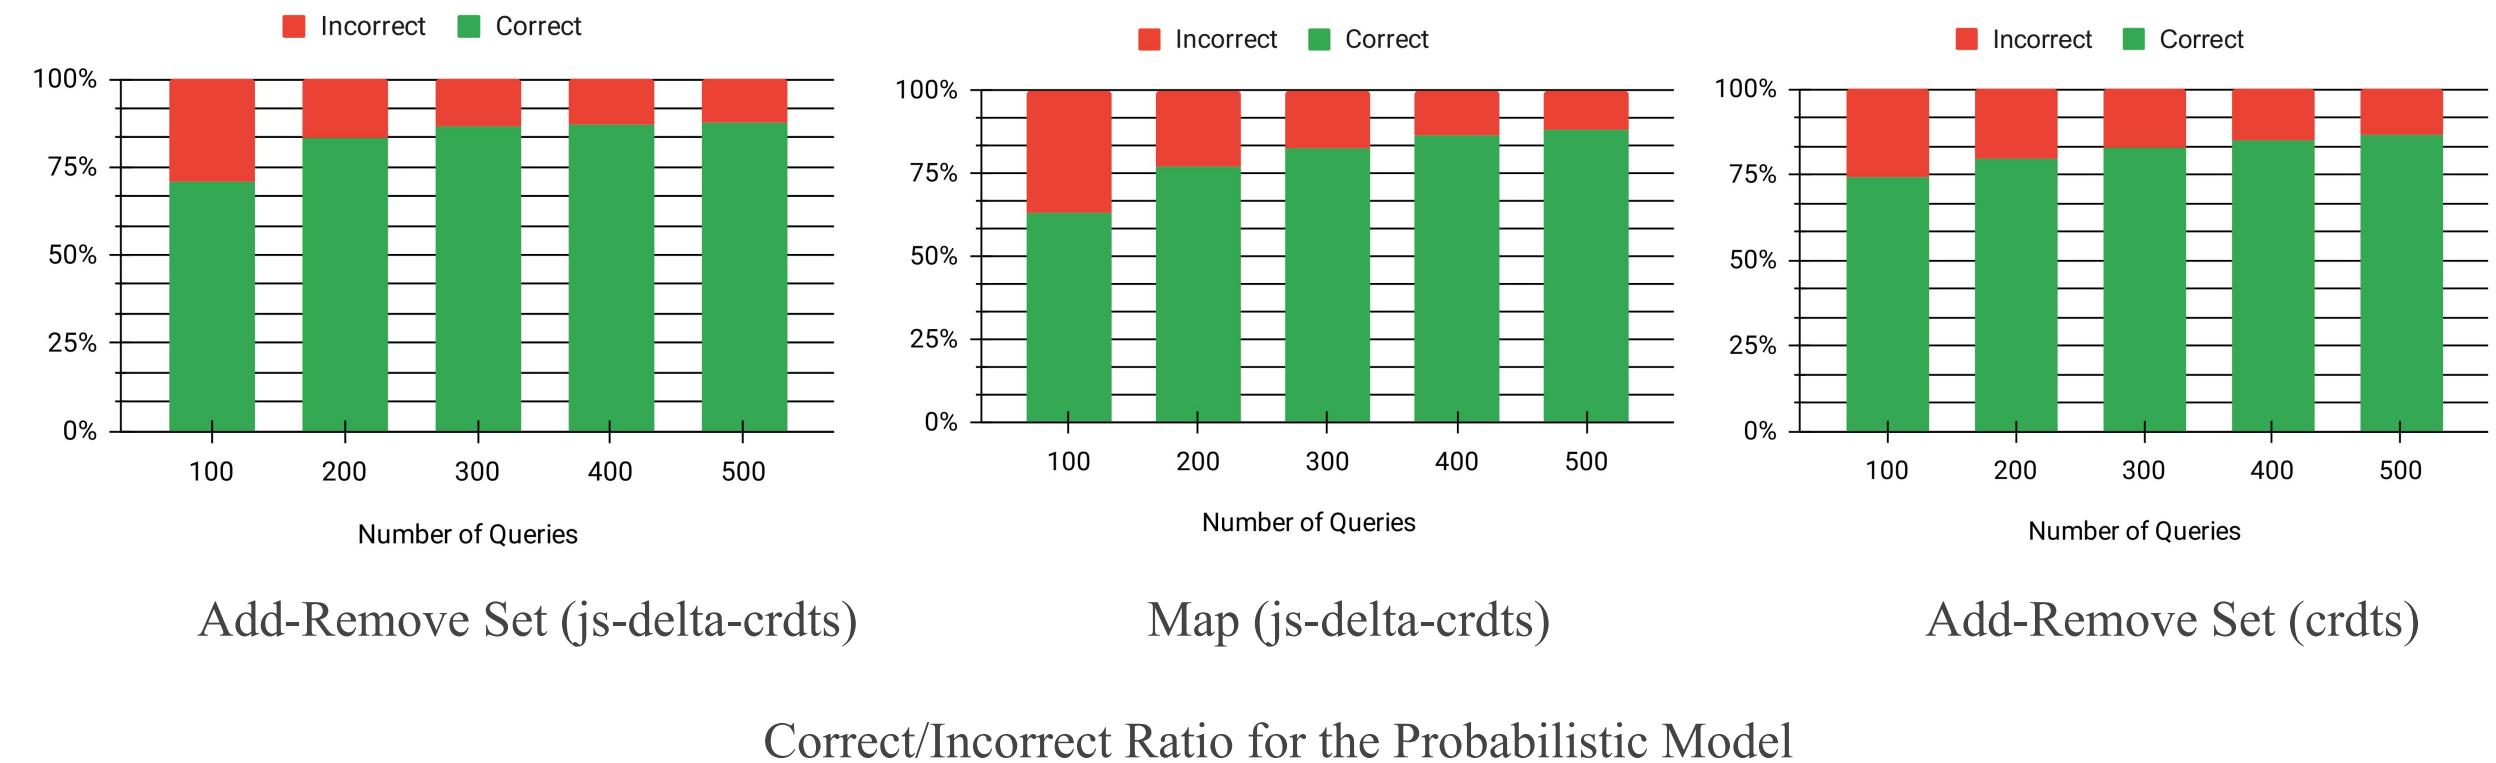
Map (js-delta-crdts)



Add-Remove Set (crdts)

Latency Comparison of Deterministic vs Probabilistic Model

RQ2: Accuracy/Performance Trade-off



Conclusion

- ❖ An automatic enhancement of CRDTs with the undo capability
- ❖ Undo functionality external to CRDT library
- ❖ Application of declarative meta-programming to CRDT
- ❖ Low performance overhead
- ❖ Ability to trade off correctness and performance