

ER- π : Exhaustive Interleaving Replay for Testing Replicated Data Library Integration

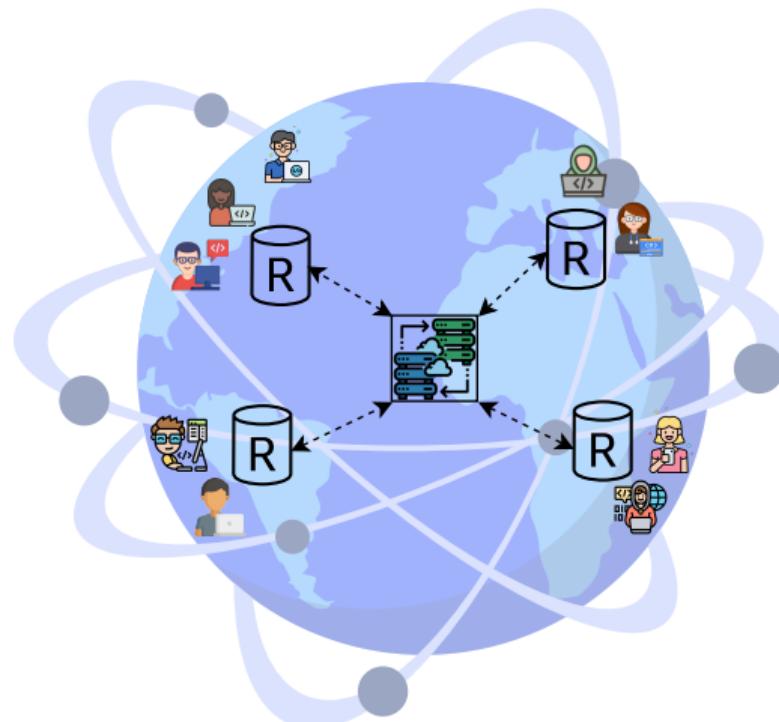
Provakar Mondal and Eli Tilevich

Software Innovations Lab, Virginia Tech
`{provakar, tilevich}@cs.vt.edu`

Table of Contents

- ① Background
- ② Problem Context
- ③ Pruning Algorithms
- ④ System Architecture
- ⑤ Evaluation
- ⑥ Limitations & Insights

Replicated Data System (RDS)



- + High availability
- + Low latency
- + Partition tolerance
- + Scalability

RDS in Practice

Domains



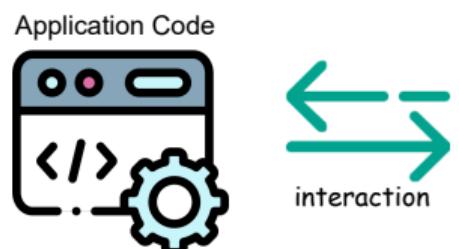
Platforms



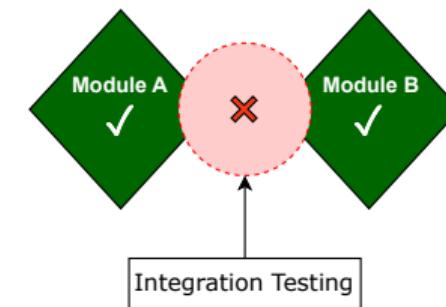
Middleware Designs

- CRDT
- MRDT
- ECRO

Integrating RDL and Testing the Integration



App ⇌ RDL Interaction



Problem Context



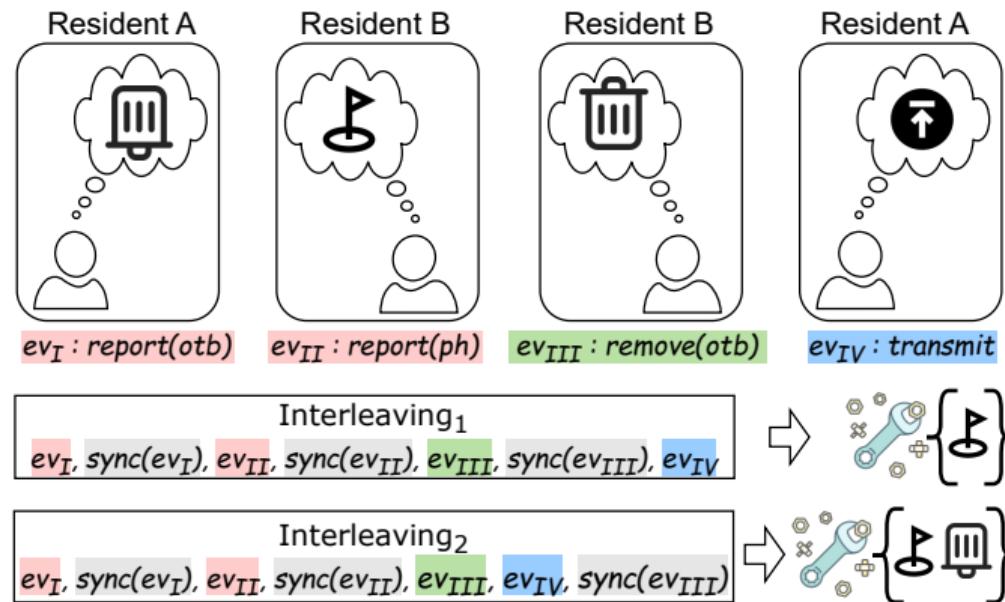
Application developers often:

- misunderstand RDL properties [1]
- hold incorrect assumptions [2]

Goals:

- ✗ Test RDL Design
- ✗ Test RDL Implementation
- ✓ Test RDL Integration

Motivating Scenario



Challenges

- Non-deterministic distributed execution
- Subtle bugs in some interleavings [3]
- Exhaustive replay for bug reproduction
- Combinatorial explosion of interleavings
- Impractical exhaustive replay

Solution Overview

ER- π —a middleware for exhaustive integration testing of RDL-based applications.

- ① Detects distributed events raised from the app-RDL interactions.
- ② Generates and persists all possible interleavings.
- ③ Reduces the problem space via *four* novel pruning algorithms.
- ④ Replays interleavings to reproduce bugs.

Pruning Algorithms

- ① Event Grouping
- ② Replica Specific
- ③ Event Independence
- ④ Failed Ops

Event Grouping

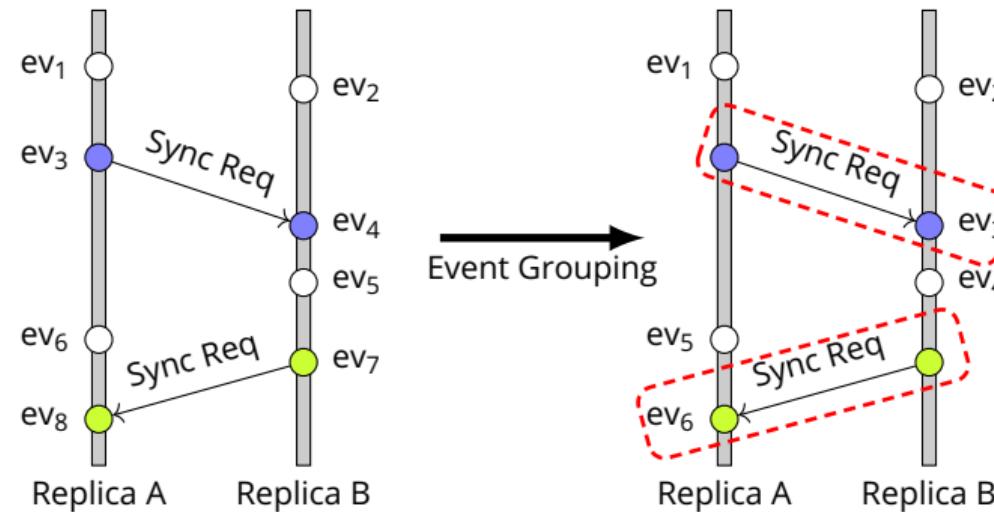


Figure: Grouping Events to Reduce Their Total #

Replica Specific

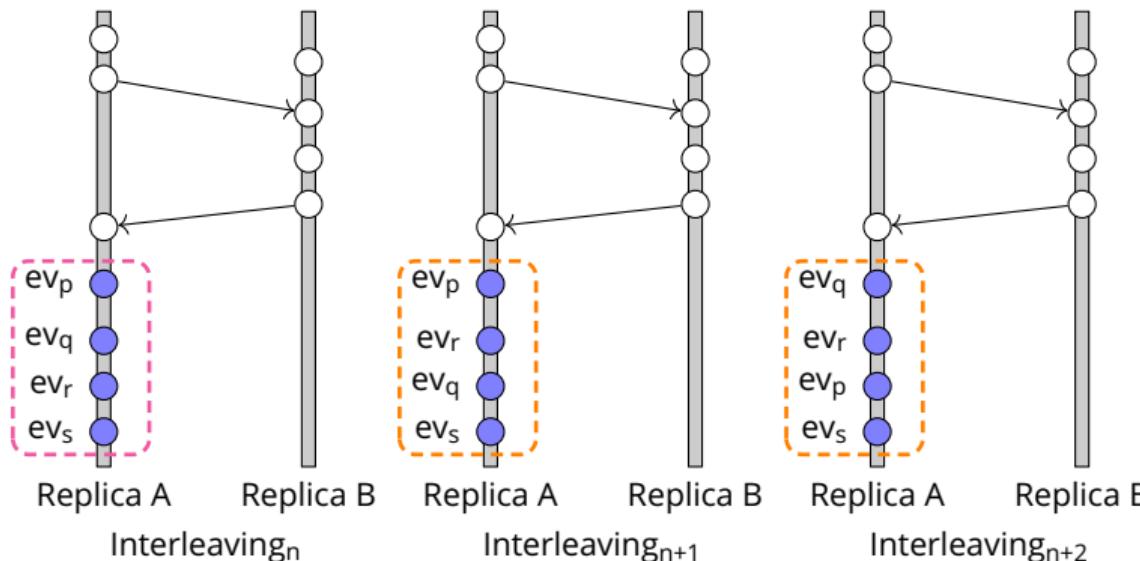
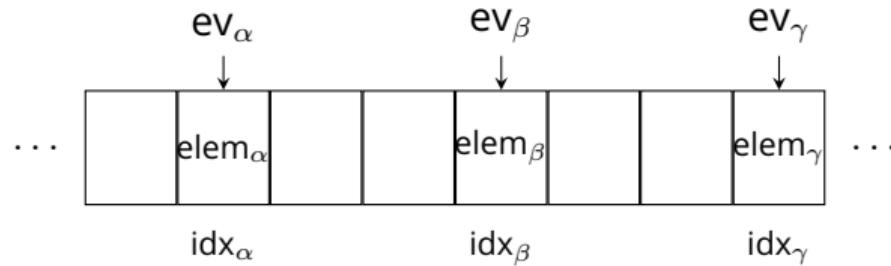


Figure: Replica B-specific Pruning

Event Independence



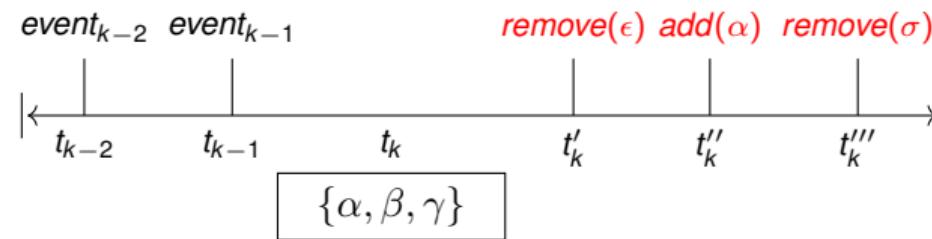
If Explored: $\{\dots \text{ev}_\alpha, \text{ev}_\beta, \text{ev}_\gamma \dots\}$

Consider Explored: $\{\dots \text{ev}_\alpha, \text{ev}_\gamma, \text{ev}_\beta \dots\}$

Consider Explored: $\{\dots \text{ev}_\beta, \text{ev}_\alpha, \text{ev}_\gamma \dots\}$

Figure: Event Independence Pruning

Failed Ops



If Explored: $\{...remove(\epsilon), add(\alpha), remove(\sigma)... \}$

Consider Explored: $\{...remove(\epsilon), remove(\sigma), add(\alpha)... \}$

Consider Explored: $\{...add(\alpha), remove(\epsilon), remove(\sigma)... \}$

Figure: Failed Ops Pruning

Motivating Example Revisited

- Sync event causally depends on the corresponding update event

① $ev_I : report(otb)$

i) $(ev_I, sync(ev_I)) \Rightarrow ev'_I$

② $sync(ev_I)$

ii) $(ev_{II}, sync(ev_{II})) \Rightarrow ev'_{II}$

③ $ev_{II} : report(ph)$

iii) $(ev_{III}, sync(ev_{III})) \Rightarrow ev'_{III}$

④ $sync(ev_{II})$

iii) $ev_{IV} : transmit$

⑤ $ev_{III} : remove(otb)$

4 events \equiv 24 Interleavings

⑥ $sync(ev_{III})$

• $\{ev_{IV}, \underbrace{\{ev'_I, ev'_{II}, ev'_{III}\}}_{3!=6}\} \rightarrow$ Problems: {} \Rightarrow 1 Interleaving

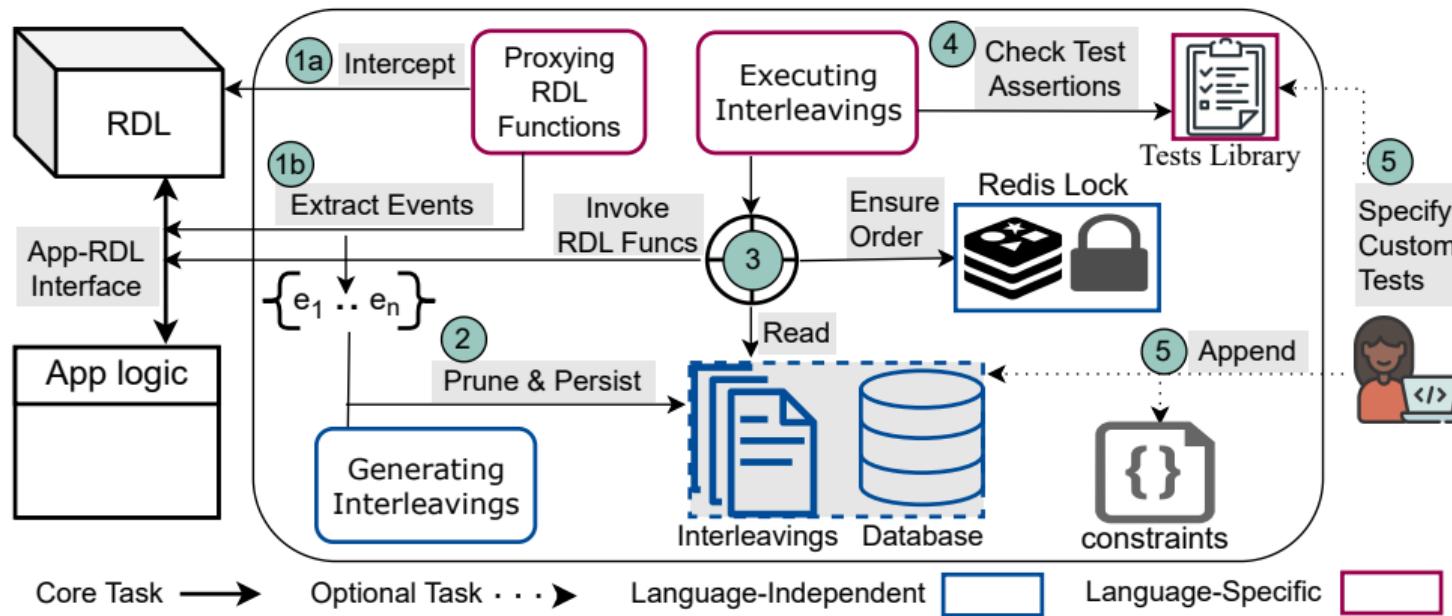
⑦ $ev_{IV} : transmit$

Total Interleavings: $24 - 5 = 19$

7 Events \equiv 5040 Interleavings

Problem Space Reduction: $\lfloor \frac{5040}{19} \rfloor = 265 \times$

System Components and Workflow



Implementation

Language-

- Specific Components: Go(≈ 300), JavaScript(≈ 280), and Java(≈ 415)
- Agnostic Components: C++($\approx 2K$)

Distributed Lock: Redis Lock [4]

Research Questions

- ✓ RQ1: Reproducing Bugs
- ✓ RQ2: Recognizing Misconceptions
- ✓ RQ3: Reducing Problem Space

RQ1: Reproducing Bugs

BugName	Issue#	#Events	Status	Reason
1. Roshi-1	18	9	closed	misconception
2. Roshi-2	11	10	closed	RDL issue
3. Roshi-3	40	21	closed	misconception
4. OrbitDB-1	513	12	open	—
5. OrbitDB-2	512	8	open	—
6. OrbitDB-3	1153	15	closed	misuse
7. OrbitDB-4	583	18	closed	misconception
8. OrbitDB-5	557	24	closed	misconception
9. ReplicaDB-1	79	10	closed	misuse
10. ReplicaDB-2	23	14	closed	misconception
11. Yorkie-1	676	17	open	—
12. Yorkie-2	663	22	closed	misconception

Table: Bug benchmarks. ‘‘#Events’’—# of interleaved events. ‘‘Status’’—if the bug is closed by library developers. ‘‘Reason’’—what causes the bug.

RQ2: Recognizing Misconceptions

- ① The underlying network ensures causal delivery [5].
- ② The order of List elements is always consistent [5].
- ③ Moving items in a List doesn't cause duplication [6].
- ④ Sequential IDs are always suitable for creating new items in a to-do list [6].
- ⑤ Multiple replicas in different regions mathematically resolve to the same state without coordination [7, 2].

RQ3: Reducing Problem Space

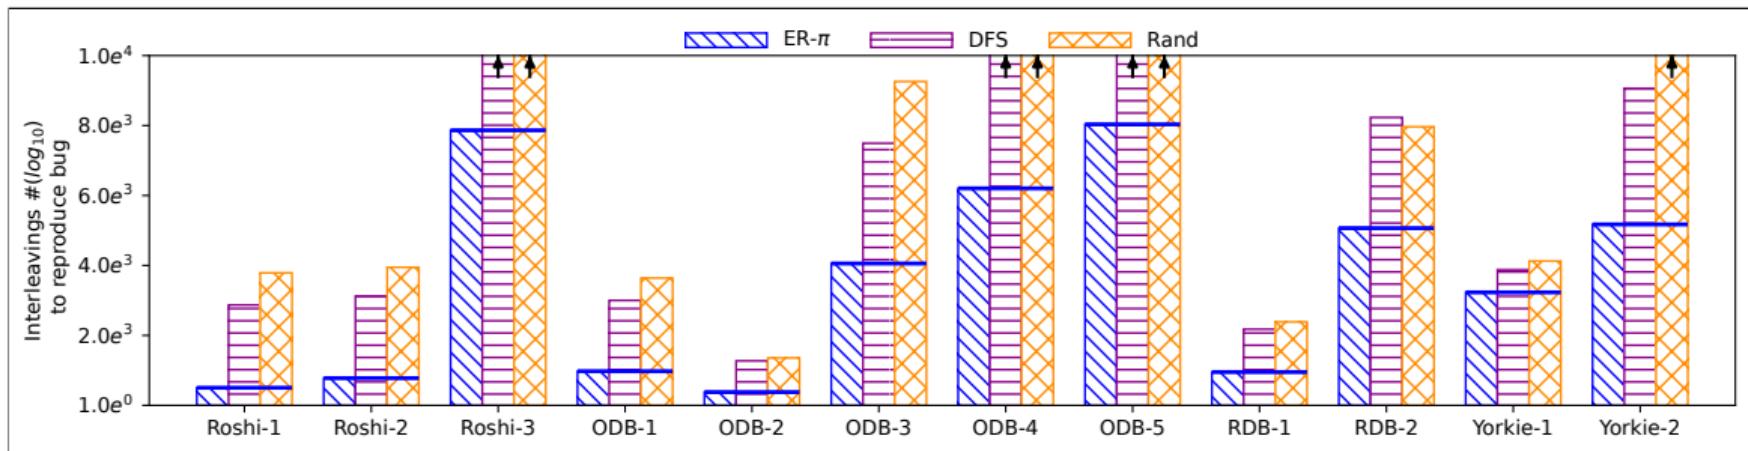


Figure: Number of Interleavings Explored to Reproduce Bug

RQ3: Reducing Problem Space

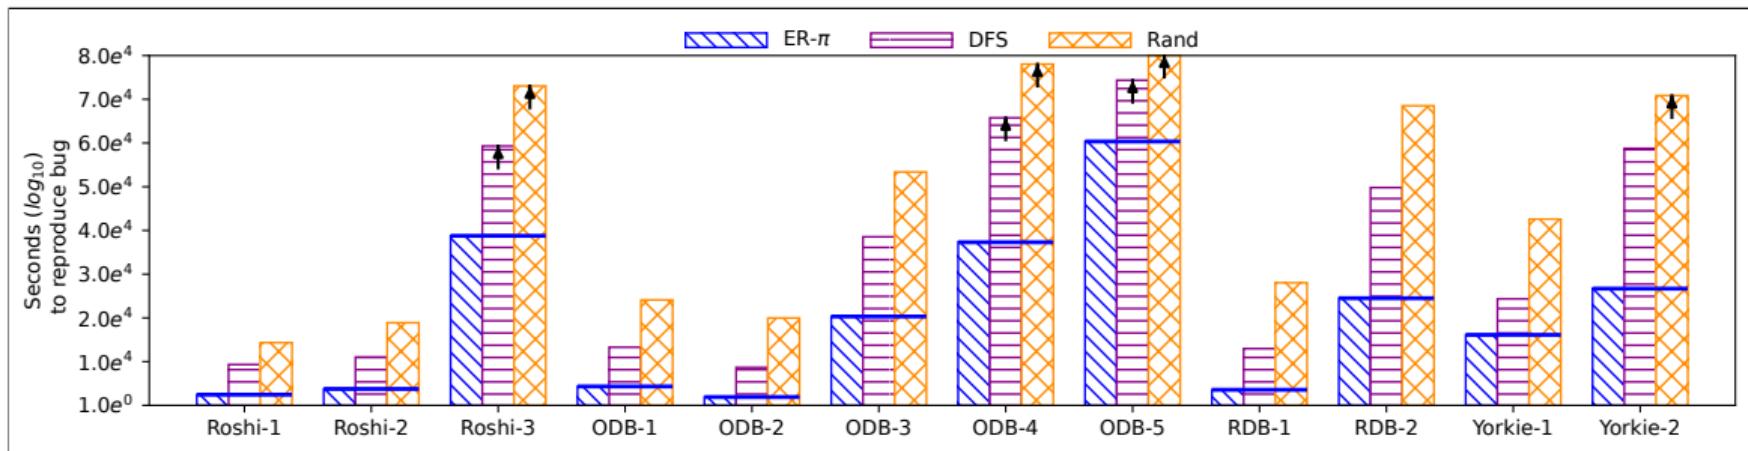


Figure: Time Required to Reproduce Bug

Limitations

- SUT comprises distinct RDL and business logic components
- Pruning algorithms assume eventual consistent RDLs
- Interfacing with RDLs relies on robust proxying support

Insights and Implications

Key Insights

- Non-deterministic interleavings can cause subtle bugs
- Misconceptions about App-RDL interactions are common
- Exhaustive interleaving replay is feasible with effective pruning
- Middleware for integration testing can benefit RDL applications

Implications

- Middleware for replaying interleavings is useful, even without pruning
- Testing RDL-based applications requires tools like ER- π
- Future RDL design should strive to minimize usage misconceptions

Summary

ER- π —a middleware for exhaustive integration testing of RDL-based applications.

- ① Detects distributed events raised from the app-RDL interactions.
- ② Generates and persists all possible interleavings.
- ③ Reduces the problem space via *four* novel pruning algorithms.
- ④ Replays interleavings to reproduce bugs.

References I

- [1] Yicheng Zhang, Matthew Weidner, and Heather Miller. "Programmer Experience When Using CRDTs to Build Collaborative Webapps: Initial Insights". In: *13th annual workshop on the intersection of HCI and PL (PLATEAU)*. Mar. 2023. DOI: [10.1184/R1/22277341.v1](https://doi.org/10.1184/R1/22277341.v1).
- [2] Shadaj Laddad et al. "Keep CALM and CRDT On". In: *Proceedings of the VLDB Endowment* 16.4 (2022), pp. 856–863.
- [3] Ohad Shacham, Mooly Sagiv, and Assaf Schuster. "Scaling Model Checking of Dataraces Using Dynamic Information". In: *Proceedings of the tenth ACM SIGPLAN symposium on Principles and practice of parallel programming*. 2005, pp. 107–118.
- [4] Redis Documentation. *Distributed Locks with Redis*.
<https://redis.io/docs/latest/develop/use/patterns/distributed-locks/>. 2023.
- [5] Yuqi Zhang et al. "MET: Model Checking-Driven Explorative Testing of CRDT Designs and Implementations". In: *arXiv preprint arXiv:2204.14129* (2022).

References II

- [6] Andrew Jeffery and Richard Mortier. "AMC: Towards Trustworthy and Explorable CRDT Applications with the Automerger Model Checker". In: *Proceedings of the 10th Workshop on Principles and Practice of Consistency for Distributed Data*. 2023, pp. 44–50.
- [7] Leena Joshi. *How to simplify distributed app development with CRDTs*. Accessed 2024. 2022.